

The `websockets` Package

Bryan W. Lewis
blewis@illposed.net

October 2, 2011

1 Introduction

The HTML 5 websocket API is a modern socket-like communication protocol for the web. The `websockets` package is a native websocket implementation for R that supports most of the draft IETF protocols in use by web browsers. The `websockets` package is especially well-suited to interaction between R and web scripting languages like Javascript. Multiple simultaneous websocket server and client connections are supported.

The library has few external dependencies and, written mostly in R, is easily portable. More significantly, `websockets` lets Javascript and other scripts embedded in web pages directly interact with R, bypassing traditional middleware layers like .NET, Java, and web servers normally used for such interaction. In some cases, websockets may make much more efficient use of networks than Ajax-like schemes for interacting with clients over web protocols.

2 Using websockets as a server, step by step

The `websockets` package includes a server function that can initiate and respond to websocket and HTTP events over a network connection (websockets are an extension of standard HTTP). All R/Websocket server applications share the following basic recipe:

1. Load the library.
2. Initialize a websocket server with `create_server`.
3. Set callback functions that will respond to desired events.
4. Service the server's socket interface with `service`, often in an event loop.

5. Shutdown the server and delete the server environment when done.

We outline the steps with examples below.

2.1 Load the library

```
library('websockets')
```

The library depends on the `caTools` and a recent version of the `digest` package. It suggests that the `RJSONIO` library be installed, as it is quite useful to have available when interacting with Javascript.

2.2 Initialize a websocket server with `create_server`

The R/Websocket service is initialized by a call to the `create_server` function. (The initialization method called `createContext` from older versions of the package is still supported.) The function takes two arguments, a network port to listen on, and an optional function closure to service standard HTTP requests (described in greater detail below). The `create_server` function returns an environment that stores data associated with the newly created server. “Callback” functions may be assigned that respond to websocket events. Here is an example that creates a websocket server on the default port of 7681:

```
server = create_server()
```

The websocket server will respond directly to any websocket client request. For convenience, the server may optionally also service basic HTML requests. For example, the basic package demo available from `demo('websockets')` serves clients the file `basic.html` located in the package installation path. However, serving HTML web pages is not the primary function of the `websockets` library—see the [Rook](#) package for that.

2.3 Set callback functions to respond to events

Clients may connect to the websocket service immediately after the server is initialized. The server may write data to or close client connections at any time. However, one must define functions to respond to incoming client events.

A server supports the following incoming events:

established: Occurs when a websocket client connection is successfully negotiated.

closed: Occurs when a client websocket connection has been closed.

receive: Occurs when data is received from a connection.

R functions may be defined to handle some, all, or none of the above event types. Such functions are termed “callbacks.”

The `set_callback` function may be used to define a callback function in the server environment returned by `create_server`. (It simply assigns the functions in that environment.) A callback function must take precisely 3 parameters that are filled in by the library with values corresponding to an event when invoking a callback function. The required parameters are:

1. **DATA**: A RAW vector that holds any incoming data associated with the event. It may be of length zero if the event does not have any data to report.
2. **WS**: The websocket client associated with the event, represented as an R list.
3. **...**: Presently unused.

The parameters can be arbitrarily named, but there must be three.

The example function below sends a message to each newly-established connection:

```
f = function(DATA, WS, ...) {  
  websocket_write("Hello there!", WS)  
}  
set_callback("established", f, server)
```

Here is an example function callback that receives data from a client connection and simply echoes it back:

```
g = function(DATA, WS, ...) {  
  websocket_write(DATA, WS)  
}  
setCallback("receive", g, server)
```

2.4 Accept requests from web clients

Javascript and other web script clients can very easily interact with the R `websockets` library directly from most browsers. The listing below presents a very basic example—see the demo scripts in the the package installation path for more complete examples.

```
<html><body>  
<script>  
socket = new WebSocket("ws://localhost:7681", "chat");  
try {  
  socket.onmessage = function got_packet(msg) {  
    document.getElementById("output").textContent = msg.data;  
  }  
}
```

```
    }  
    catch(ex) {document.getElementById("output").textContent = "Error: " + ex;}  
</script>  
<div id="output"> SOCKET DATA APPEARS HERE </div>  
</body></html>
```

Note: The `websockets` package presently ignores the sub-protocol (“chat” in the above example). Future versions of the package may require specific sub-protocols.

2.5 Service the socket interface with `service`

Websocket events are placed in a queue. The `service` function processes events in the queue on a first-come, first-served basis. The `service` function processes each event by invoking the appropriate callback function. It returns after a configurable time out if there are no events to service. Events may be processed indefinitely by evaluating the `service` function in a loop, for example:

```
while(TRUE)  
{  
  service(server)  
}
```

The `service` function timeout value prevents the R session from spinning and consuming lots of CPU time. See the `service` help page for more information.

2.6 Sending data to clients

The `websocket_write` and `websocket_broadcast` functions may be used to send data to clients. The `websocket_broadcast` function emulates a broadcast by sending data in a loop to all connected websocket clients associated with the specified server.

The `websocket_write` function may be used at any time to send data to a specific websocket client. Each websocket server environment returned by the `create_server` function maintains a list of connected client sockets in the variable `client_sockets`. Each client socket is in turn represented by an R list. The following example assumes that the `server` environment has been initialized and contains a single client:

```
websocket_write("Hello", server$client_sockets[[1]])
```

Note the use of the double bracket indexing operator to select a single list element from the `client_sockets` list.

2.7 Close the server when done

Servers should be closed when done as follows:

```
websocket_close(server)
```

2.8 HTTP convenience functions

The `websockets` package includes two convenience function closures for servicing basic HTTP requests: `static_file_service` and `static_text_service`. The functions take either a file name or text string that contains an HTML web page, respectively, and issue a well-formed HTTP 200 response to the requesting client. They are intended to be used in the `webpage` argument to the `create_server` function. The following example defines a basic web page in a string variable:

```
content='<html><body>
  <script>
    socket = new WebSocket("ws://localhost:7681", "chat");
    try {
      socket.onmessage = function got_packet(msg) {
        document.getElementById("output").textContent = msg.data;
      }
    }
    catch(ex) {document.getElementById("output").textContent = "Error: " + ex;}
  </script>
  <div id="output"> SOCKET DATA APPEARS HERE </div>
</body></html>
'

server = create_server(webpage=static_text_service(content))
```

The web page text will be issued to any client making an HTTP GET request. To serve content from files instead, use the `static_file_service` function. Additionally, the `static_file_service` function checks to see if the file has been updated and always uses a fresh version.

Note that both convenience functions mostly ignore the GET RESOURCE and all other GET request parameters. They always only return the specified HTML content. POST requests are always ignored by the websocket server. Users are free to define their own function closures to use instead, which may be more full-featured. Use the existing functions as a guide. If you really need a full-featured HTTP service, we recommend using the [Rook](#) package instead.

3 Using websockets as a client, step by step

WRITE ME!

4 Tips and miscellaneous notes

We present a few more advanced and other miscellaneous notes in this section.

4.1 Binary data

Binary data is supported by IETF websocket protocol versions greater than 00. The `websockets` package supports the older 00 protocol with ASCII-only data, as well as binary data transfers with newer clients.

JSON is probably a good choice to use when interacting with Javascript and the data size is not too large. The suggested `RJSONIO` package helps map many native R objects to JSON and vice versa, greatly facilitating interaction between R and Javascript. But, JSON data is transferred as characters, which may incur performance and in some cases numeric issues.